

Tech Talks

# Chasing Races

## Deterministic Simulation Testing in XTDB

James Henderson, XTDB Head of Engineering

2025-12-05

# Coming up:

- The **value** and **cost** of a test
- What we test in XTDB, and how
- DST: the theory
- What this looks like in Kotlin

# The Value and Cost of a Test



# Value $\propto$ ?

- **False negative** rate
- **Probability** of a bug in the code-under-test
- **Impact** of a bug in the code-under-test
- Documentation value
- others...?



# Cost $\propto$ ?

- **Time** to write it (obviously)
- **False positive** rate
- **Maintenance** cost
- Execution time
- Setup
- others...?



# What we test on XTDB, and how?

- ‘Highly general’ code: example-based **unit tests**
  - e.g. bitemporal resolution, what to compact
- Complex outputs: **regression tests**
  - e.g. SQL planner, indexer
- Combinatorial explosions: **property tests**
  - e.g. polymorphic Arrow vectors
- Concurrent/distributed code: **deterministic simulation tests** ← you are here
- Everything else: mostly example-based **integration tests**
  - using the in-memory, throwaway node setup - highly recommend

# Aside: Allen intervals

- Useful in our unit tests + property tests

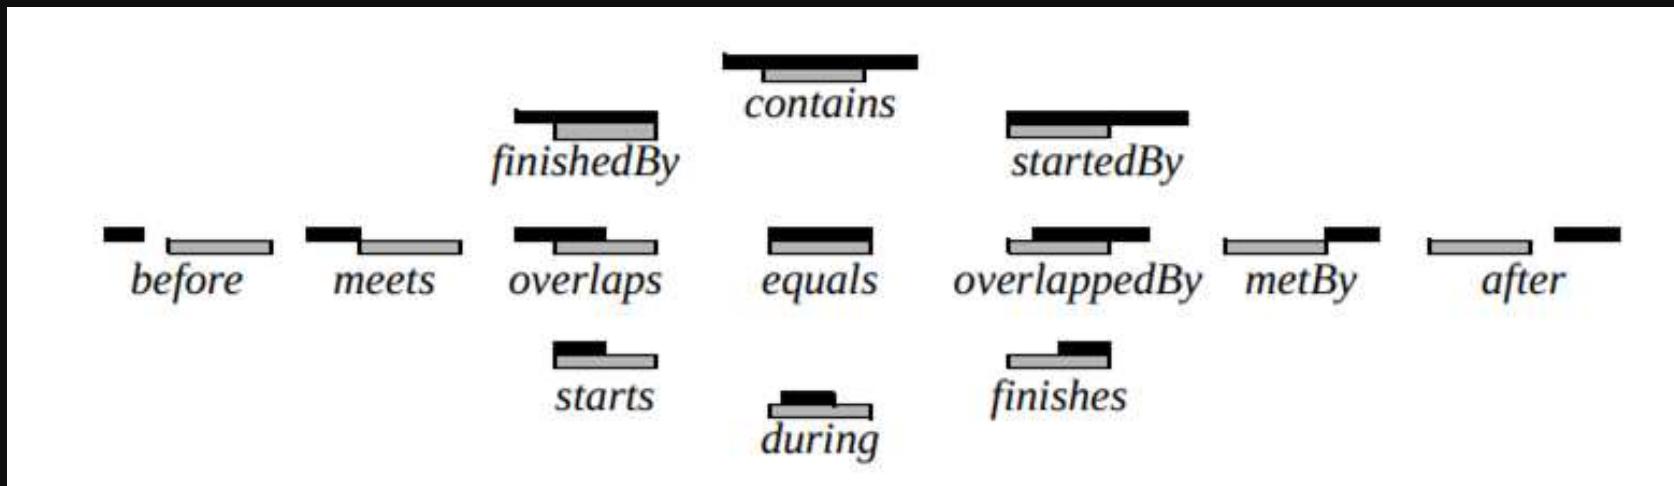


image: "Object-Relational Indexing for General Interval Relationships", Kriegel et al.

# Before you DST:

- DST is *relatively* complex (compared to unit/integration/property)
- Like any risk, **avoid** before you **mitigate**:
  - Immutability by default
  - Avoid sharing mutable state between threads
  - Separate pure code from side-effects
  - “Communicating sequential processes” (CSP), actors

... a.k.a. ‘Functional Programming Best Practices’

```
sealed interface Event
class FetchReq(val key: Key, val res: CompletableDeferred<Value>) : Event
class FetchDone(val key: Key, val value: Value) : Event

val fetchCh = Channel<Event>()

init {
    scope.launch {
        val state = State(/* ... */)

        for (event in fetchCh) {
            when (event) {
                is FetchReq -> /* ... */
                is FetchDone -> /* ... */
            }
        }
    }
}

suspend fun get(key: Key): Value {
    val res = CompletableDeferred<Value>()
    fetchCh.send(FetchReq(key, res))
    return res.await()
}
```

# Deterministic Simulation Testing (DST)

Problem:

- Concurrent code is non-deterministic - timing dependent
- Vast number of potential ‘interleavings’
- Anything you’d normally do to isolate the issue (e.g. logging, debugger) often makes the issue go away

Solution: just run all the possible interleavings, see if anything breaks!

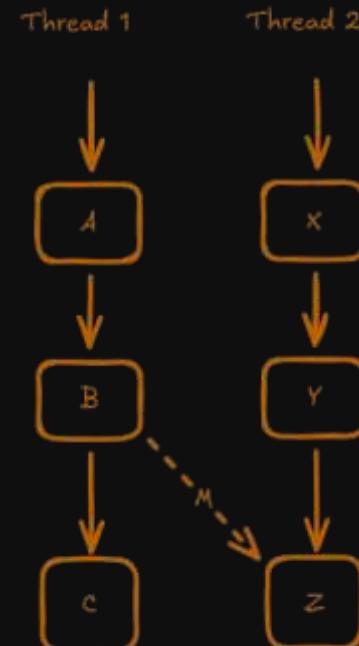
Thank you for coming to my TED talk

# 'happens before'

1. If events A and B are in the same thread, and A comes before B in that thread, then  $A \rightarrow B$
2. If A is the sending of a message by one process, and B is the receiving of that same message by another process, then  $A \rightarrow B$
3. If  $A \rightarrow B$  and  $B \rightarrow C$ , then  $A \rightarrow C$  ('transitivity')

'happens before' is only a **partial order**.

If  $\neg(A \rightarrow B)$  and  $\neg(B \rightarrow A)$  then A and B are **concurrent**, and can happen in either order.



*"Time, Clocks, and the Ordering of Events in a Distributed System"* (Lamport, 1978)

# DST: Trying all the total orderings

Aims:

- Choose an ordering, run it, then test properties of the system
- If an ordering fails any of its properties, allow exact re-running
  - e.g. with a debugger attached, logging turned up.
- Low false positive and false negative rates
  - Test the ‘real code’ as much as possible
  - ... so need to minimise performance overhead
- Keep code-under-test simple
  - i.e. introducing DST shouldn’t require complex changes to the underlying code

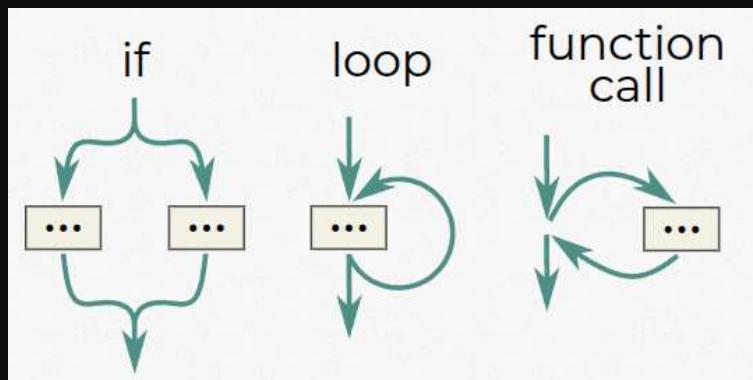
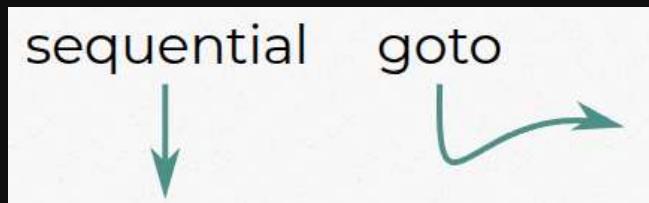
# Coroutines and Continuations:

```
fun calculateSum(a: Int, b: Int): Int = a + b
```

```
suspend fun calculateSum(client: HttpClient, a: Int, b: Int): Int =  
    client.get("http://example.com/sum?a=$a&b=$b").body()
```

# Aside: ‘structured’ concurrency

a.k.a. “Go Statement Considered Harmful”

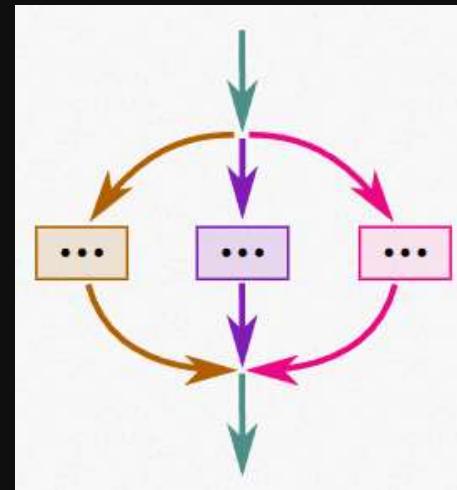
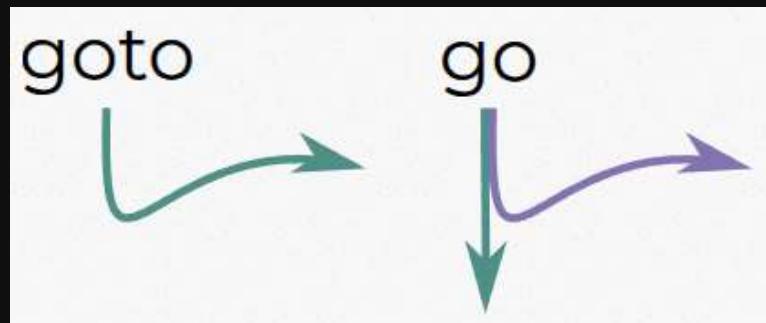


*Left: A traditional `goto`. Right: A domesticated `goto`, as seen in C, C#, Golang, etc. The inability to cross function boundaries means it can still pee on your shoes, but it probably won't rip your face off.*



# Aside: ‘structured’ concurrency

a.k.a. “Go Statement Considered Harmful”



# Coroutines and Continuations:

```
suspend fun calculateSum(client: HttpClient, a: Int, b: Int): Int =  
    client.get("http://example.com/sum?a=$a&b=$b").body()
```

```
fun calculateSum(  
    client: HttpClient, a: Int, b: Int,  
    cont: Continuation<Int>  
): Unit {  
    client.get(  
        "http://example.com/sum?a=$a&b=$b",  
        Continuation { res -> cont.resume(res.body()) }  
    )  
}
```

# CoroutineDispatcher

```
public abstract class CoroutineDispatcher {  
  
    /**  
     * Requests execution of a runnable [block].  
     *  
     * The dispatcher guarantees that [block] will eventually execute,  
     * typically by dispatching it to a thread pool, using a dedicated thread,  
     * or just executing the block in place.  
     *  
     * This method should guarantee that the given [block] will be eventually invoked,  
     * otherwise the system may reach a deadlock state and never leave it.  
     */  
    public abstract fun dispatch(context: CoroutineContext, block: Runnable)  
}
```

# CoroutineDispatcher

```
class DeterministicDispatcher(private val rand: Random) : CoroutineDispatcher() {  
  
    private val jobs = mutableSetOf<Runnable>()  
    private var running = false  
  
    override fun dispatch(context: CoroutineContext, block: Runnable) {  
        jobs.add(block)  
  
        if (!running) {  
            running = true  
            while (true) {  
                val job = jobs.randomOrNull(rand) ?: break  
                jobs.remove(job)  
                job.run()  
            }  
            running = false  
        }  
    }  
}
```

# Isolating the observable side-effects

Goal:

every time there's a observable side-effect,  
give the Dispatcher chance to re-order.

# Isolating the observable side-effects

Goal:

every time there's a observable side-effect,  
give the Dispatcher chance to re-order.

```
interface Driver : AutoCloseable {  
    suspend fun executeJob(job: Job): TriesAdded  
    suspend fun appendMessage(triesAdded: TriesAdded): Log.MessageMetadata  
}
```

- Mock Driver uses `yield()`'s to *ensure* suspension

# Spot the bug (Apache Arrow Java)

```
/**  
 * Increment the ledger's reference count  
 * by the given amount.  
 */  
  
fun retain(increment: Int) {  
    require(increment > 0) { "retain must increment" }  
  
    val originalReferenceCount = bu...  
  
    require(originalReferenceCount > 0)  
}
```

apache / arrow-java

<> Code ⚡ Issues 401 Pull requests 26 Discussions Actions Security Insights

Calling `retain` on a closed ArrowBuf revives it, into an undefined state #906

[Open](#)

 jarohen opened last month

Describe the bug, including details regarding any error messages, version, and platform.

`#Test`  
`fun testArrowBufRetainBug():`  
 `RootAllocator().use { al ->`  
 `val buf = al.buffer(10)`  
 `buf.close()`  
  
 `// correctly throws`  
 `assertThrows<IllegalArgumentException> { buf.referenceManager.retain() }`  
 `assertEquals(0, buf.refCount()) // note, it's 1`  
  
 `// doesn't throw, ref-count is 1 - so the caller assumes they've successfully taken a reference`  
 `// but the underlying memory has already been deallocated and re-used`  
 `assertThrows<IllegalArgumentException> { buf.referenceManager.retain() }`  
 }

Caused by the `getAndAdd` in `BufferLedger.retain(int)` - this is what leaves the ref-count positive, so on the next call, this doesn't fail.

Some kind of compareAndSet instead, perhaps?

Cheers,  
James

Q Type ↗ to search

# Yeah, ok - but why's that a race?

- Two threads trying to increment the ref-count:
  - Buffer ref-count → 0, fair game for cleanup.
  - T1 tries, (correctly) gets the exception - evicts it from our cache.
  - T2 tries at that exact moment - still in cache, but ***no exception*** (ref-count = 1, thanks T1)
  - T2 reads freed memory - gonna have a *bad time*.
- Check the ref-count first, and then try to increase it?
  - That's also a race

# Now for a test:

```
@RepeatableSimulationTest
fun `deterministic concurrent fetch of same path-slice`(iteration: Int) = runTest {
    MemoryCache(/* ... */).use { cache ->
        val path = Path.of("...")

        // Launch multiple concurrent fetches of the same path-slice
        val deferreds = async(dispatcher) {
            (1..5).map { i ->
                async(dispatcher) {
                    cache.get(path).use { buf ->
                        assertEquals(/* ... */)
                        yield()
                        buf[0]
                    }
                }
            }
        }.await()

        val results = deferreds.awaitAll()

        assertEquals(/* all results are the same */)
    }
}
```

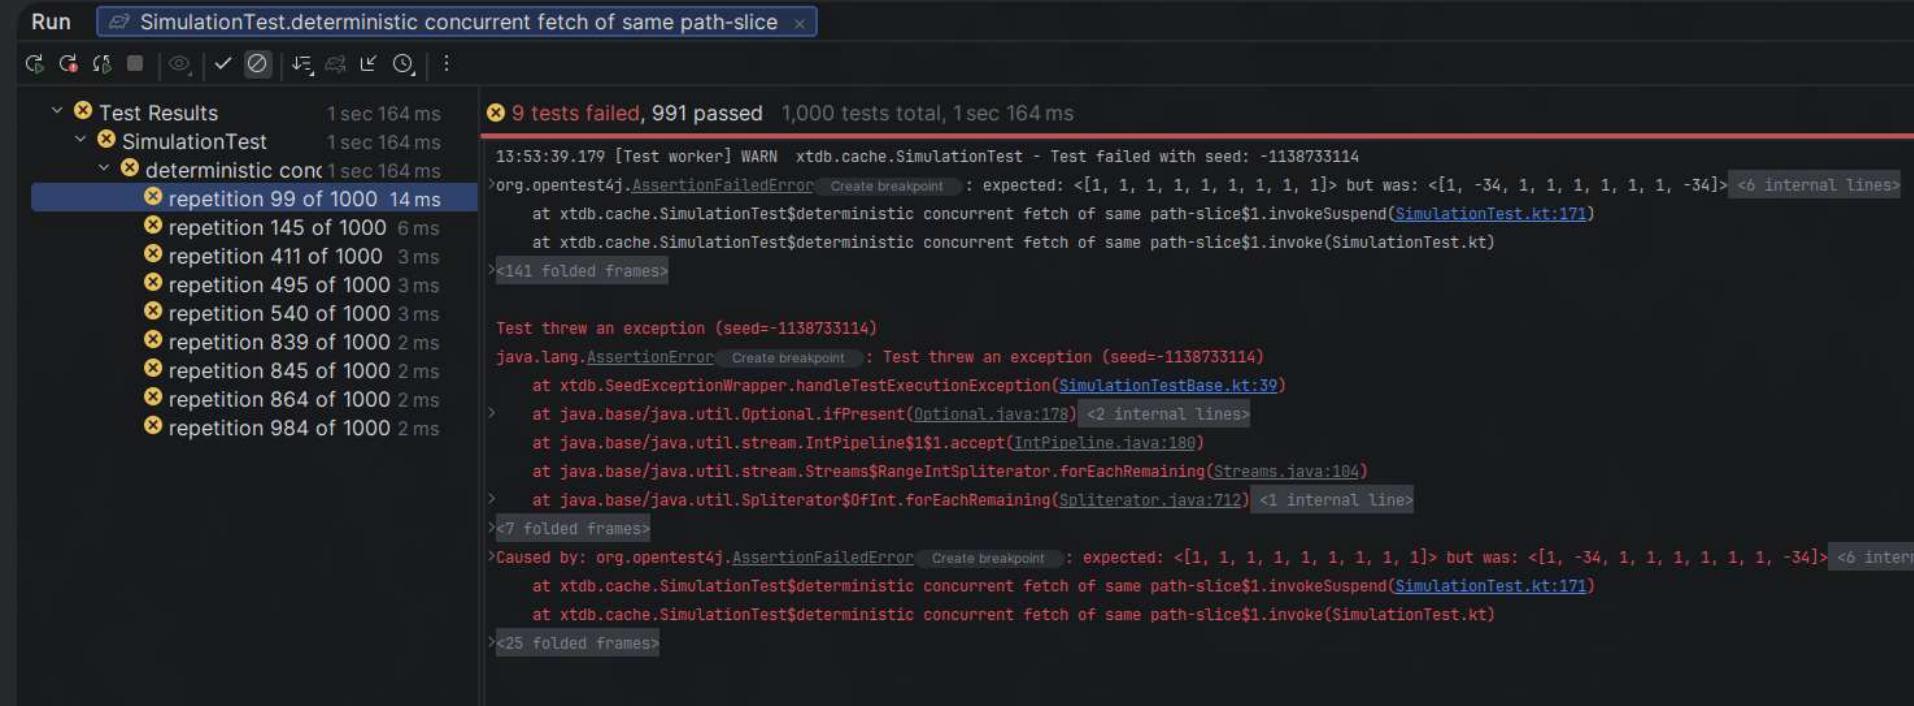
```
class DeterministicDispatcher(private val rand: Random) : CoroutineDispatcher() {

    private val jobs = mutableSetOf<Runnable>()
    private var running = false

    override fun dispatch(context: CoroutineContext, block: Runnable) {
        jobs.add(block)

        if (!running) {
            running = true
            while (true) {
                val job = jobs.randomOrNull(rand) ?: break
                jobs.remove(job)
                job.run()
            }
            running = false
        }
    }
}
```







Run

SimulationTest.deterministic concurrent fetch of same path-slice

# The fix

```
/**  
 * Increment the ledger's reference count for associated underlying memory chunk  
 * by the given amount.  
 */  
fun retain(increment: Int) {  
    require(increment > 0) { "retain($increment) argument is not positive" }  
  
    while (true) {  
        val currentRefCount = bufRefCnt.get()  
        check(currentRefCount > 0)  
        if (bufRefCnt.compareAndSet(currentRefCount, currentRefCount + increment))  
            return  
    }  
}
```



Run

SimulationTest.deterministic concurrent fetch of same path-slice



- ✓ Test Results 1sec 229 ms
- ✓ SimulationTest 1sec 229 ms
- ✓ deterministic concul 1sec 229 ms
- ✓ repetition 1 of 1000 185 ms
- ✓ repetition 2 of 1000 4 ms
- ✓ repetition 3 of 1000 4 ms
- ✓ repetition 4 of 1000 7 ms
- ✓ repetition 5 of 1000 5 ms
- ✓ repetition 6 of 1000 4 ms
- ✓ repetition 7 of 1000 5 ms
- ✓ repetition 8 of 1000 4 ms
- ✓ repetition 9 of 1000 4 ms
- ✓ repetition 10 of 1000 5 ms
- ✓ repetition 11 of 1000 4 ms
- ✓ repetition 12 of 1000 4 ms
- ✓ repetition 13 of 1000 4 ms
- ✓ repetition 14 of 1000 4 ms
- ✓ repetition 15 of 1000 4 ms
- ✓ repetition 16 of 1000 3 ms
- ✓ repetition 17 of 1000 3 ms
- ✓ repetition 18 of 1000 3 ms
- ✓ repetition 19 of 1000 2 ms
- ✓ repetition 20 of 1000 2 ms
- ✓ repetition 21 of 1000 3 ms
- ✓ repetition 22 of 1000 4 ms
- ✓ repetition 23 of 1000 4 ms
- ✓ repetition 24 of 1000 4 ms
- ✓ repetition 25 of 1000 4 ms
- ✓ repetition 26 of 1000 3 ms
- ✓ repetition 27 of 1000 2 ms
- ✓ repetition 28 of 1000 2 ms
- ✓ repetition 29 of 1000 2 ms
- ✓ repetition 30 of 1000 2 ms
- ✓ repetition 31 of 1000 2 ms
- ✓ repetition 32 of 1000 2 ms
- ✓ repetition 33 of 1000 3 ms

✓ 1,000 tests passed 1,000 tests total, 1sec 229 ms



# Closing thoughts

- FP Best Practices™ are still as true as ever
  - Immutability by default
  - Separate data and behaviour
  - Separate pure from side-effecting code, test them in isolation → simpler tests
- “Communicating sequential processes” / “actors” have been a boon
  - If you’re going to have mutable state, definitely don’t have *shared* mutable state
- ... but if you do have shared mutable state, then DST

# Thank you Merry Christmas!

•••

@jms (JUXT)

@jhenderson (Grid)

@jarohen (everywhere else)

